Bayesian Data Analysis

Vanda Inácio & Ken Newman

University of Edinburgh

Semester 2, 2017/2018

Regression

Linear regression: general context

- ↪ Linear regression is, perhaps, the most widely used statistical modelling tool.
- ↪ It addresses the following question: How does a quantity of primary interest, y, vary as (depend upon) another quantity, or set of quantities, $\mathbf{x}$?
- ↪ The quantity y is called the response or outcome variable. Some people simply refer to it as the dependent variable.
- ↪ The variable(s) $\mathbf{x} = (x_1, x_2, \ldots, x_p)$ are called explanatory variables, covariates or simply independent variables.
- ↪ For example, y might be the final mark in the Bayesian data analysis class. Variable x_1 might be the number of lectures attended, x_2 the amount of time spent studying, and so on. Regression allows us to test the effect of, say, the number of lectures attended, but accounting for the time spent studying.

Regression

Linear regression: examples

Some examples (of varying quality!) in the media:

Does chocolate make you clever?

By Charlotte Pritchard
BBC News

Eating more chocolate improves a nation's chances of producing Nobel Prize winners - or at least that's what a recent study appears to suggest. But how much chocolate do Nobel laureates

In today's Magazine

12 March 2012 Last updated at 20:17

Share

Red meat increases death, cancer and heart risk, says study

COMMENTS (729)

A diet high in red meat can shorten life expectancy, according to researchers at Harvard Medical School.

The study of more than 120,000 people suggested red meat increased the risk of death from cancer and heart problems.

Substituting red meat with fish, chicken or nuts lowered the risks, the authors said.

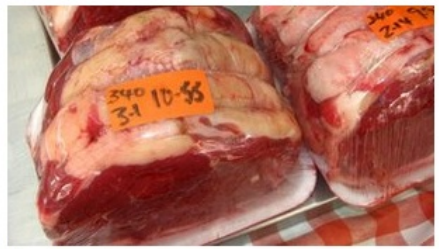

Experts advise to choose leaner cuts of red meat

The British Heart Foundation said red meat could still be eaten as part of a balanced diet.

The researchers analysed data from 37,698 men between 1986 and 2008 and 83,644 women between 1980 and 2008.

They said that during the study period, adding an extra portion of unprocessed red meat to someone's daily diet would increase the risk of death by 13%, of fatal cardiovascular disease by 18% and of cancer mortality by 10%. The figures for processed meat were higher, 20% for

Related Stories

Red meat study: Risks 'very clear'

Cut red meat to lower cancer risk

How much red meat should we eat?

Regression

Linear regression

News › Society › Cancer

No clear evidence that mobile phone radiation damages health

Largest review yet of published research finds that so far there has been 'no indication' of increased risk

Ian Sample, science correspondent
theguardian.com, Thursday 26 April 2012 06.30 BST
Jump to comments (76)

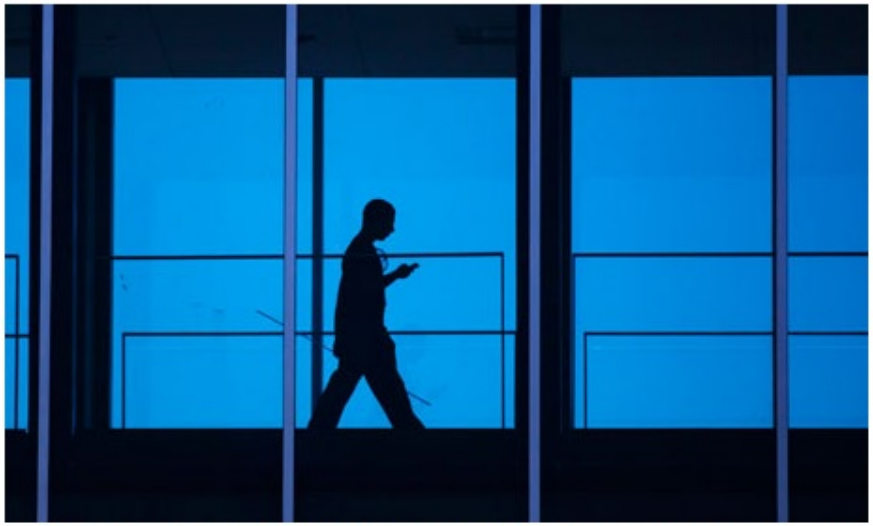

A report by a Health Protection Agency advisory group has found 'no indication' of an increased risk to health from mobile phone use. Photograph: Mark Blinch/Reuters

How a short nap can raise the risk of diabetes: Study finds people who have a siesta are more likely to have high blood pressure and high cholesterol

- **Napping for more than 30 minutes at a time can raise the risk of diabetes, according to a new study**
- **It can also increase likelihood of high blood pressure and high cholesterol**

By PAT HAGAN

PUBLISHED: 01:04, 21 September 2013 | **UPDATED:** 10:34, 21 September 2013

Share Tweet +1 Share 598 shares 102 View comments

They were much favoured by Margaret Thatcher, Albert Einstein and Winston Churchill.

But while afternoon naps may revitalise tired brains, they can also increase the risk of diabetes, according to new research.

A study of more than 27,000 people in China – where taking a post-lunch snooze is very popular – shows napping for more than 30 minutes at a time can raise the chances of developing type two diabetes.

Researchers found men and women taking 40 winks were also more likely to have high blood pressure and raised cholesterol levels compared to those who stayed awake through the day.

Regression

Linear regression: examples

↪ The two main aims of regression analysis are

1. **Estimation**: the estimation of relationships between the explanatory and response variables with the object of both identifying and quantifying this relationship.

2. **Prediction**: the prediction of new values of y from values of $x_1, \ldots x_p$.

Regression

Linear regression: formulation

- We will start by considering regression models of the form

$$y_i \overset{\text{ind}}{\sim} \text{N}(\mu_i, \sigma^2), \qquad i = 1, \ldots, n \tag{1}$$

where

$$\mu_i = \beta_0 + \beta_1 x_{i1} + \ldots + \beta_p x_{ip}.$$

- The variables $\beta_0, \beta_1, \ldots \beta_p$ are an unknown set of regression coefficients (to be estimated from the data).

- Equation (1) is equivalent to the following relationship

$$y_i = \beta_0 + \beta_1 x_{i1} + \ldots + \beta_p x_{ip} + \epsilon_i, \qquad i = 1, \ldots, n$$

where $\epsilon_i \overset{\text{iid}}{\sim} N(0, \sigma^2)$.

Regression

Linear regression: formulation

↪ This is because, since the x_{ij} are known, we can write

$$\begin{aligned}\mathbb{E}\left(y_i \mid x_{i1}, \ldots x_{ip}\right) &= \mathbb{E}\left(\beta_0 + \beta_1 x_{i1} + \ldots \beta_p x_{ip} + \epsilon_i\right) \\ &= \beta_0 + \ldots + \beta_p x_{ip} \\ &= \mu_i\end{aligned}$$

$$\begin{aligned}\text{var}(y_i \mid x_{1i}, \ldots x_{pi}) &= \text{var}(\beta_0 + \beta_1 x_{i1} + \ldots + \beta_p x_{ip} + \epsilon_i) \\ &= \text{var}(\epsilon_i) \\ &= \sigma^2\end{aligned}$$

↪ In short, $y_i = \mathbf{x}_i^T \boldsymbol{\beta} + \varepsilon_i$, where $x_i = (1, x_{i1}, \ldots, x_{ip})^T$ and $\boldsymbol{\beta} = (\beta_0, \beta_1, \ldots, \beta_p)^T$.

Regression

Linear regression: formulation

↪ In matrix notation, we have

$$\mathbf{y} = \mathbf{x}\beta + \varepsilon,$$

where $\mathbf{y} = (y_1, \ldots, y_n)^T$, $\varepsilon = (\varepsilon_1, \ldots, \varepsilon_n)^T$,

$$\mathbf{x} = \begin{bmatrix} 1 & x_{11} & x_{12} & \ldots & x_{1p} \\ 1 & x_{21} & x_{22} & \ldots & x_{2p} \\ \vdots & \vdots & \vdots & \vdots & \vdots \\ 1 & x_{n1} & x_{n2} & \ldots & x_{np} \end{bmatrix},$$

and $\varepsilon \sim \mathrm{N}(\mathbf{0}, \sigma^2 I_n)$.

Regression

Linear regression: formulation

- There are $p + 2$ unknown parameters to be estimated:
  - $\beta = (\beta_0, \beta_1, \ldots \beta_p)^T$
  - σ^2
- The data $(y_i, x_{i1}, x_{i2}, \ldots x_{ip})$ for $i \in \{1, \ldots n\}$ are used to obtain estimates.
- Under the classical setting, the estimator for $\boldsymbol{\beta}$ is given by

$$\hat{\boldsymbol{\beta}} = (\mathbf{x}^T\mathbf{x})^{-1}\mathbf{x}^T\mathbf{y},$$

assuming that $(\mathbf{x}^T\mathbf{x})$ is invertible. Since we are assuming normal errors, both maximum likelihood and ordinary least squares give rise to this estimator.

Regression

Linear regression: Bayesian formulation

- We will now turn our attention to Bayesian estimation of $\boldsymbol{\beta}$ and σ^2.
- The Bayesian linear model is given by

$$
\begin{aligned}
y_i \mid \mu_i, \sigma^2, \mathbf{x}_i &\overset{\text{ind}}{\sim} \text{N}(\mu_i, \sigma^2), \qquad i = 1, \ldots, n, \\
\mu_i &= \mathbf{x}_i^T \boldsymbol{\beta}, \\
\boldsymbol{\beta}, \sigma^2 &\sim p(\boldsymbol{\beta}, \sigma^2).
\end{aligned}
$$

- $p(\boldsymbol{\beta}, \sigma^2)$ is the joint prior on the parameters. There are several approaches to model the joint prior distribution.
- All inference proceeds from $p(\boldsymbol{\beta}, \sigma^2 \mid \mathbf{y})$.
- The full conditionals $p(\beta_j \mid \tau, \mathbf{y})$ are not difficult, but tedious, to obtain. The full conditional distribution $p(\tau \mid \boldsymbol{\beta}, \mathbf{y})$, where $\tau = \sigma^{-2}$, is trivially obtained.

Regression

Linear regression: Bayesian formulation

↪ In normal regression models, the simplest approach is to assume that all parameters are a priori independent having the structure

$$p(\boldsymbol{\beta}, \tau) = \prod_{j=0}^{p} p(\beta_j) p(\tau), \qquad \tau = \sigma^{-2},$$

$$\beta_j \sim \mathrm{N}(\mu_{\beta_j}, \sigma^2_{\beta_j}), \quad j = 0, \ldots, p,$$

$$\tau \sim \mathrm{Gamma}(a, b).$$

↪ In this prior setup, again, we have substituted the variance σ^2 by the corresponding precision parameter τ in order to make it compatible to the BUGS notation.

Regression

Linear regression: Bayesian formulation

- When no information is available, a usual choice for the prior mean is the zero value ($\mu_{\beta_j} = 0$).
- This prior choice centers our prior beliefs around zero, which corresponds to the assumption of no effect of x_j on y.
- In this way, we express our prior doubts about the effect of x_j on y, prompting Spiegelhalter et al. (2004, pp. 90, 158-160) to call this a "sceptical" prior.
- The prior variance $\sigma^2_{\beta_j}$ of the effect β_j is set equal to a large value (e.g., 10^4) to represent high uncertainty or prior ignorance.
- A very popular prior for τ, especially between BUGS users, is to use equally low prior parameter values, e.g., $a = b = 0.1$, or $a = b = 0.01$. Both prior configurations lead to a peak in 0^+. Their use is controversial. For a good discussion on this topic, see

  `http://www.stat.columbia.edu/~gelman/research/published/taumain.pdf`

Regression

Linear regression: Bayesian formulation

- ↪ Each regression coefficient pertains to the effect of explanatory variable x_j on the expectation of the response variable y adjusted for the remaining covariates.
- ↪ To ascertain whether the effect of x_j is important for the prediction or description of y, we initially focus on examining whether the posterior distribution of β_j is scattered around zero (or not).
- ↪ Posterior distributions far away from the zero value will indicate an important contribution of x_j on the prediction of the response variable.
- ↪ Although formal Bayesian hypothesis testing is not based on simply examining the posterior distribution and their credible intervals, such analysis can offer a first and reliable tool for tracing important variables.

Regression

Linear regression: Bayesian formulation

↪ Further, the magnitude of the effect of variable x_j on y is given by the posterior distribution of β_j ($j = 1, \ldots, p$) since

$$\begin{aligned}
&\mathbb{E}(y \mid x_1, \ldots, x_{j-1}, x_j = x+1, x_{j+1}, \ldots, x_p) - \mathbb{E}(y \mid x_1, \ldots, x_{j-1}, x_j = x, x_{j+1}, \ldots, x_p) \\
=&\beta_0 + \beta_j(x+1) + \sum_{k \neq j, k=1}^{p} \beta_k x_k - \beta_0 - \beta_j x - \sum_{k \neq j, k=1}^{p} \beta_k x_k \\
=&\beta_j
\end{aligned}$$

↪ Hence the posterior mean or median of β_j will correspond to the corresponding posterior measures of the expected change of the response variable y.

↪ Hence, an increase of one unit of x_j, given that the remaining covariates will remain stable, induces an a posteriori average change on the expectation of y equal to the posterior mean of β_j.

↪ In short, β_j is the expected change (increase or decrease) when x_j increases by one unit and the rest of the covariates remain unchanged.

Regression

Linear regression: Bayesian formulation

- ↪ Concerning the intercept parameter β_0, its interpretation corresponds to the expected value of the response variable y when the observed values of all covariates are equal to zero.
- ↪ Frequently such combination lies outside the range of the observed covariate values.
- ↪ In many cases, direct interpretation of β_0 does not lead to realistic and sensible interpretation.
- ↪ An alternative is to center around zero all explanatory variables x_j, by subtracting their sample mean. In this case, the constant β_0 represents the expected value of y when all covariates are equal to its sample means, representing in this way the expected response y for an "average" or "typical" subject according to our sample.
- ↪ Note that this does not change the meaning of β_j, $j = 1, \ldots, p$.
- ↪ Centering (or centering and scaling) has some further advantages: it usually helps the mixing of the MCMC chains and also facilitates prior specification.

Regression

Linear regression

↪ Let us look at an example. The following dataset (the `mtcars` dataset available on `R`) gives the fuel consumption of cars, together with three aspects of car construction and performance.

	mpg	Rear axle ratio	Weight (lb/1000)	1/4 mile time
Mazda RX4	21.0	3.90	2.620	16.46
Mazda RX4 Wag	21.0	3.90	2.875	17.02
Datsun 710	22.8	3.85	2.320	18.61
Hornet 4 Drive	21.4	3.08	3.215	19.44
Hornet Sportabout	18.7	3.15	3.440	17.02
Valiant	18.1	2.76	3.460	20.22
Duster 360	14.3	3.21	3.570	15.84
⋮	⋮	⋮	⋮	⋮

↪ Based on these data we are interested in estimating the fuel consumption of a car based on the rear axle ratio, the weight, and the 1/4 mile time.

Regression

Linear regression

- We will fit a Bayesian multiple regression model to the `mtcars` dataset.
- The model we will fit is

$$y_i = \beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \beta_3 x_{i3} + \varepsilon_i, \quad \varepsilon_i \overset{\text{iid}}{\sim} N(0, \sigma^2), \quad i = 1, \ldots 32,$$

  where

  - y_i is the miles driven per gallon for the i-th car.
  - x_{i1} is the rear axle ratio for the i-th car.
  - x_{i2} is the weight of the i-th car (lb/1000).
  - x_{i3} is the 1/4 mile time for the i-th car.
- We will consider $\beta_j \sim \text{N}(0, 1000)$ (variance of 1000 or precision of 0.001) and $\tau \sim \text{Gamma}(0.1, 0.1)$.

Regression

Linear regression

↪ We start by fitting a frequentist regression model (using ordinary least squares).

```
data(mtcars)
help(mtcars)

#creating a new dataframe with only mpg,drat,wt, and qsec as variables
vars=names(mtcars)%in%c("cyl", "disp", "hp", "vs","am","gear","carb") #variables to exclude
mtcars1=mtcars[!vars]

fit=lm(mpg~drat+wt+qsec,data=mtcars1)
summary(fit)

Call:
lm(formula = mpg ~ drat + wt + qsec, data = mtcars1)

Residuals:
    Min      1Q  Median      3Q     Max
-4.1152 -1.8273 -0.2696  1.0502  5.5010

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  11.3945     8.0689   1.412  0.16892
drat          1.6561     1.2269   1.350  0.18789
wt           -4.3978     0.6781  -6.485 5.01e-07 ***
qsec          0.9462     0.2616   3.616  0.00116 **

Residual standard error: 2.56 on 28 degrees of freedom
Multiple R-squared:  0.837,Adjusted R-squared:  0.8196
F-statistic: 47.93 on 3 and 28 DF,  p-value: 3.723e-11
```

Regression

Linear regression

↪ We now implement the model using the package `rjags`. We need to specify the likelihood of our normal multiple linear regression model.

```
require(rjags)
n=nrow(mtcars1)
#covariates and response
y=mtcars1$mpg; drat=mtcars1$drat; wt=mtcars1$wt; qsec=mtcars1$qsec

model_string <- "model{

# Likelihood
for(i in 1:n){
y[i]~dnorm(mu[i],tau)
mu[i]=beta[1]+beta[2]*drat[i]+beta[3]*wt[i]+beta[4]*qsec[i]
}

# Prior for beta
for(j in 1:4){
beta[j]~dnorm(mu0,tau0)
}
tau0=1/sigma02

# Prior for the precision
tau~dgamma(a, b)

# Compute the variance
sigma2=1/tau
}"
```

Regression

Linear regression

↪ We will now pass all information needed and compile the model.

```
#hyperparameters for the betas and tau
mu0=0; sigma02=1000; a=0.1; b=0.1

# list with data and hyperparameters
data=list(y=y,drat=drat,wt=wt,qsec=qsec,n=n,mu0=mu0,sigma02=sigma02,a=a,b=b)

#Compile the model
model=jags.model(textConnection(model_string),n.chains=1,data=data)

Compiling model graph
   Resolving undeclared variables
   Allocating nodes
Graph information:
   Observed stochastic nodes: 32
   Unobserved stochastic nodes: 5
   Total graph size: 254

Initializing model
```

↪ Note that we have 5 unobserved stochastic nodes (=variables): β_0, β_1, β_2, β_3, and σ^2. The 32 observed stochastic nodes correspond to the 32 y's.

Regression

Linear regression

↪ We will do a burn-in of 100 000 iterations, run the model for 500 000 iterations, and do a thinning of 50.

```
update(model,100000,progress.bar="none")
res=coda.samples(model,variable.names=c("beta","sigma2"),n.iter=500000,thin=50,progress.bar=
summary(res)

Iterations = 100050:6e+05
Thinning interval = 50
Number of chains = 1
Sample size per chain = 10000

1. Empirical mean and standard deviation for each variable,
   plus standard error of the mean:

           Mean     SD Naive SE Time-series SE
beta[1] 10.3694 8.0796 0.080796       0.234179
beta[2]  1.7771 1.2427 0.012427       0.029145
beta[3] -4.3354 0.6834 0.006834       0.015359
beta[4]  0.9677 0.2675 0.002675       0.005713
sigma2   6.9785 1.9893 0.019893       0.020370

2. Quantiles for each variable:

           2.5%     25%     50%    75%  97.5%
beta[1] -5.0983  5.0167 10.3319 15.617 26.349
beta[2] -0.7206  0.9583  1.7916  2.622  4.166
beta[3] -5.7024 -4.7828 -4.3340 -3.885 -2.995
beta[4]  0.4487  0.7894  0.9673  1.149  1.493
sigma2   4.1601  5.5733  6.6201  8.016 11.729
```

Regression

Linear regression

- ↪ Note that the above results might slightly change when you run the model on your computer (stochastic nature of the MCMC).
- ↪ The chains for the regression coefficients were highly correlated, and that is why we have thinned them.
- ↪ Effective sample size for each parameter (β_0, β_1, β_2, β_3, and σ^2) are: 1190, 1817, 1979, 2192, and 9536, respectively.
- ↪ All convergence diagnostics seem ok!
- ↪ As an exercise, try centering and scaling each covariate (i.e., for each covariate subtract its mean and divide by its standard deviation) and look at autocorrelations (no thinning). Does it help?

Regression

Linear regression

↪ Under the frequentist analysis we have

Parameter	Point Estimate (95% confidence interval)
β_0	$11.395\ (-5.134, 27.922)$
β_1	$1.750\ (-0.857, 4.169)$
β_2	$-4.347\ (-5.787, -3.009)$
β_3	$0.946\ (0.410, 1.482)$
σ^2	6.554

↪ For the Bayesian version results are as follows

Parameter	Posterior mean (95% credible interval)
β_0	$10.369\ (-5.098, 26.349)$
β_1	$1.777\ (-0.721, 4.166)$
β_2	$-4.335\ (-5.702, -2.995)$
β_3	$0.968\ (0.449, 1.493)$
σ^2	$6.978\ (4.160, 11.729)$

↪ Results are pretty similar! However, interpretation of the intervals are completely different.

Regression

Linear regression

- The regression model described earlier can be defined in JAGS using vectors and matrices instead of scalar nodes, by using the function `inprod`.

```
x=cbind(rep(1,n),drat,wt,qsec)

# Likelihood
for(i in 1:n){
y[i]~dnorm(mu[i],tau)
mu[i]=inprod(beta[],x[i,])
}

# Prior for beta
for(j in 1:4){
beta[j]~dnorm(mu0,tau0)
}
tau0=1/sigma02

# Prior for the inverse variance
tau~dgamma(a, b)

# Compute the variance
sigma2=1/tau
}"

mu0=0; sigma02=1000; a=0.1; b=0.1
data=list(y=y,x=x,n=n,mu0=mu0,sigma02=sigma02,a=a,b=b)
```

- Everything proceeds then similarly.

Regression

Linear regression

- We can also use multivariate normal priors for the regression coefficients $\boldsymbol{\beta}_{p+1} \sim \text{N}(\boldsymbol{\mu}_{\boldsymbol{\beta}}, \boldsymbol{\Sigma}_{\boldsymbol{\beta}})$, we just need to make use of the function `dmnorm`.

- Note that the prior we have used before it is a particular case (when the off-diagonal elements are all zero).

```
model_string <- "model{

# Likelihood
for(i in 1:n){
y[i]~dnorm(mu[i],inv.var)
mu[i]=inprod(beta[],x[i,])
}

beta~dmnorm(mu.beta,tau.beta)

# Prior for the inverse variance
inv.var~dgamma(a, b)

# Compute the variance
sigma2=1/inv.var
}"

#this coincides with the previous prior, as off diagonal entries are zero
mu.beta=rep(0,4); tau.beta=diag(0.0001,4); a=0.001; b=0.001
data=list(y=y,x=x,n=n,mu.beta=mu.beta,tau.beta=tau.beta,a=a,b=b)
```

Regression

Linear regression: model checking

↪ The conclusions of any (Bayesian) analysis are conditional on the appropriateness of the assumed statistical model, so we need to be satisfied that our assumptions are a reasonable approximation to reality, even though we do not generally believe any model is actually "true" (Lunn et al. 2012).

↪ In defining the multiple regression model we have assumed

1. **Normality**: the error terms $\varepsilon_1, \varepsilon_2, \ldots, \varepsilon_n$ are normally distributed, with mean 0 and variance σ^2.
2. **Independence**: the error terms $\varepsilon_1, \varepsilon_2, \ldots, \varepsilon_n$ are independent of one another.
3. **Constant variability**: the value of σ^2 does not depend on the values x_{ij} (this is known as homoskedasticity).
4. **Linearity**: the mean μ_i is a linear function of $x_{i1}, \ldots, x_{ip}$.

↪ All of these assumptions can be (informally) checked using the studentised residuals.

Regression

Linear regression: model checking

↪ To check the aforementioned listed assumptions, studentised residuals are often used in frequentist analysis

$$\widehat{\varepsilon}_i = \frac{y_i - \hat{y}_i}{\hat{\sigma}\sqrt{1 - h_{ii}}}, \qquad i = 1, \ldots, n$$

where $\hat{y}_i = \sum_{j=0}^{p} \widehat{\beta}_j x_{ij}$ and h_{ii} is the i-th diagonal entry of the 'hat' matrix $H = \mathbf{x}(\mathbf{x}^T\mathbf{x})^{-1}\mathbf{x}^T$.

↪ Under the assumption of normality of the error terms ε_i,

$$\{\widehat{\varepsilon}_i\} \dot{\sim} \mathrm{N}(0, 1)$$

Note: this is an approximation. Although the errors $\{\varepsilon_i\}$ should be independent under our model assumptions, the studentised residuals $\{\widehat{\varepsilon}_i\}$ are not necessarily independent, so tests of Normality are not applied rigorously - it is usually enough to inspect graphs, etc.

Regression

Linear regression: model checking

- In the Bayesian context, it can be read in Lunn et al. (2012, p.141)

  "If we want to create a single-valued residual rather than a random quantity, alternative possibilities include using a single draw $\theta^{(j)} = (\boldsymbol{\beta}^{(j)}, \sigma^{(j)})$*, plugging in the posterior means of* θ*, or using the posterior mean residual."*

- Here, we opt for the last option (the posterior mean of the residual), that is, for each draw $(\boldsymbol{\beta}^{(j)}, \sigma^{(j)})$ from the posterior distribution, we compute the studentised residuals and from this ensemble the posterior mean (or median if preferred) is computed.

- The same option is taken by Wang et al. (Bayesian Regression Modeling with INLA, 2018, CRC/Chapman & Hall).

Regression

Linear regression: model checking

- ↪ To assess the normality assumption, a Q-Q plot of the Bayesian studentised residuals must be constructed.
- ↪ In what regards the independence of the error terms assumption, the usual way to assess it is to plot the Bayesian studentised residuals against the order in which observations were gathered. Any pattern suggest that the errors $\{\varepsilon_i\}_{i=1}^{n}$ are not independent.
- ↪ To check the assumptions of linearity and constant variance we can plot the Bayesian studentised residuals against the posterior mean (or median) of the fitted values. This graph should look random, with no obvious patterns, if these two assumptions are valid.
- ↪ Interpreting these plots is, of course, kind of subjective...

Regression

Linear regression: model checking

↪ For the mtcars datastet, the following figures were produced

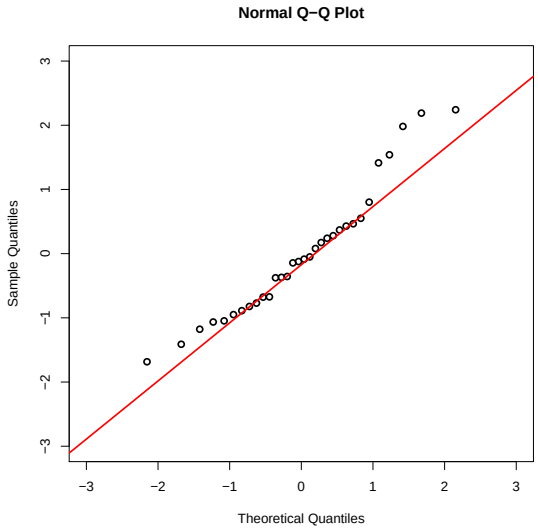

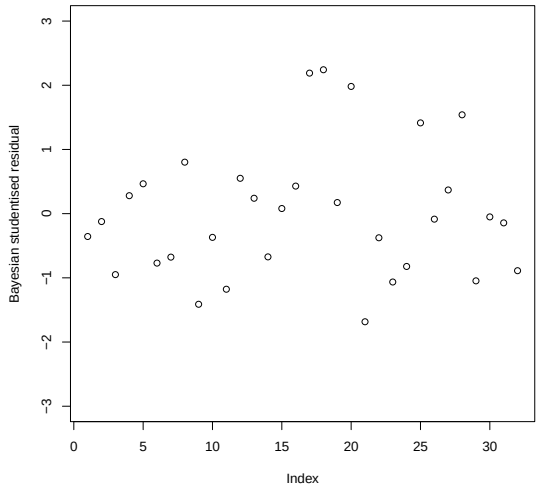

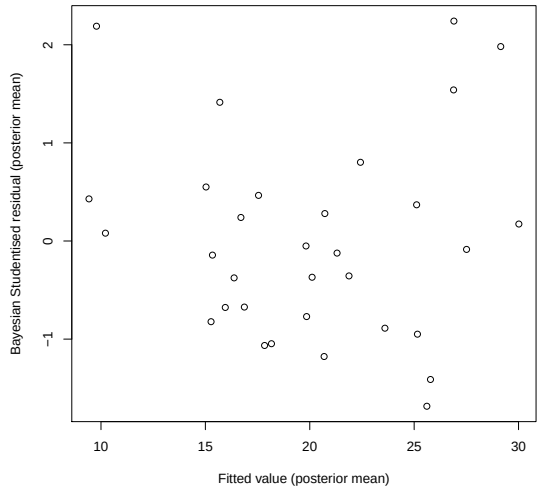

↪ All the assumptions seem to be reasonably satisfied.

Regression

Linear regression: model checking

- Predictive checks are also very popular in Bayesian statistics. Basically, for an overall assessment of model fit we can generate replicate data sets (of the same size of the original data) from the posterior predictive distribution, and compare to the real data using specific test quantities (Gelman et al. 2013, chapter 6).

- Let $\mathbf{y}_{\text{rep}}$ denotes the replicated datasets. The distribution of $\mathbf{y}_{\text{rep}}$ is the posterior predictive distribution

$$p(\mathbf{y}_{\text{rep}} \mid \mathbf{y}) = \int p(\mathbf{y}_{\text{rep}} \mid \theta) p(\theta \mid \mathbf{y}) \mathrm{d}\theta.$$

- We do not evaluate this integral but sample from $p(y_{\text{rep}} \mid \mathbf{y})$:

  1. Draw $\theta^{(j)}$ from $p(\theta \mid \mathbf{y})$, $j = 1, \ldots, M$.
  2. Draw $\mathbf{y}_{\text{rep}}^{(j)} \sim p(\mathbf{y}_{\text{rep}} \mid \theta^{(j)})$, $j = 1, \ldots, M$,

  where M is the number of MCMC iterations.

- The chosen test statistics are then computed for each replicated dataset and 'contrasted' to the one based on the observed data.

Regression

Linear regression: model checking

- We distinguish between the replicated data $\mathbf{y}_{\text{rep}}$ and the predictive outcomes $\tilde{\mathbf{y}}$.
- The variable $\tilde{\mathbf{y}}$ is any *future* observable value of the outcome. For example, in a linear regression model $\tilde{\mathbf{y}}$ can have its own set of explanatory variables $\tilde{\mathbf{x}}$.
- On the other hand, $\mathbf{y}_{\text{rep}}$ must have the same explanatory variables $\mathbf{x}$ as those used in the model for the observed data $\mathbf{y}$. In this sense, $\mathbf{y}_{\text{rep}}$ is similar to "predicting the observed data".

Regression

Linear regression: model checking

↪ A very good and concise discussion about this topic can be found on the paper

`https://arxiv.org/pdf/1709.01449.pdf`

↪ As the authors of the aforementioned paper state: "*Posterior predictive checking makes use of the data twice, once for the fitting and once for the checking. Therefore it is a good idea to choose statistics that are orthogonal to the model parameters. If the test statistic is related to one of the model parameters, e.g., if the mean statistic is used for a Gaussian model with a location parameter, the posterior predictive checks may be less able to detect conflicts between the data and the model.*"

↪ For the `mtcars` example, since we are using a location-scale normal model, let us pick up as test statistics the minimum, the maximum, the median, and the skewness.

Regression

Linear regression: model checking

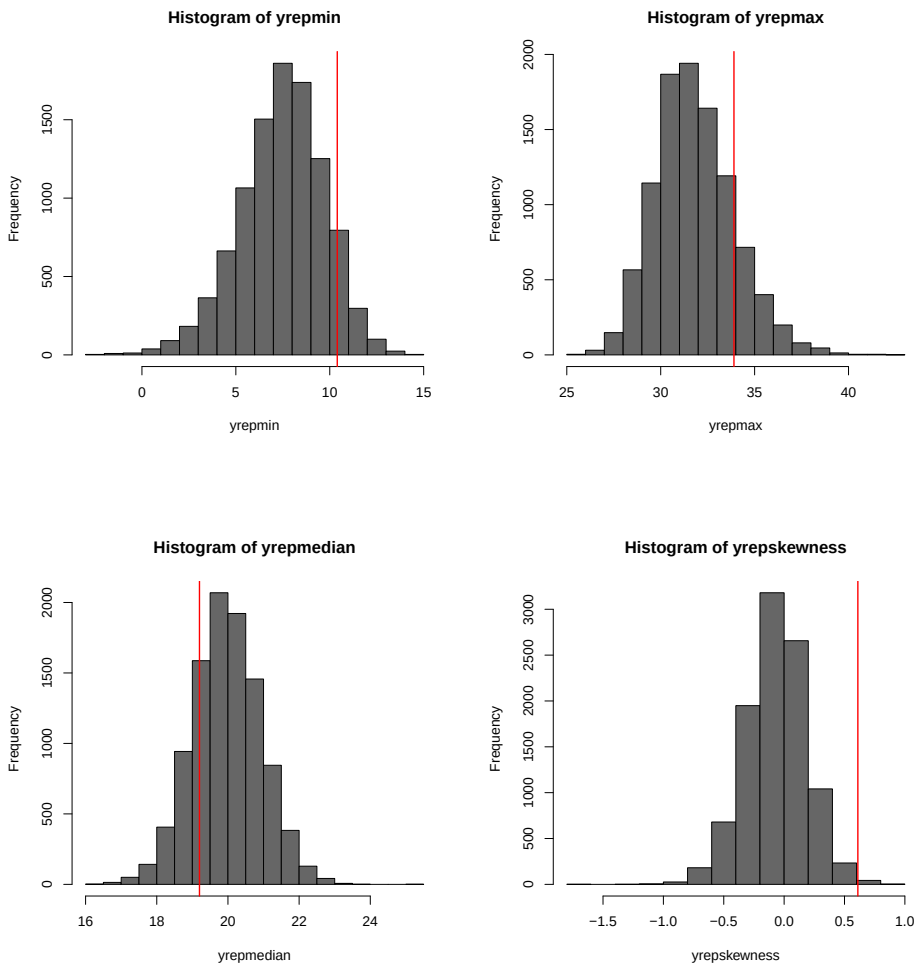

↪ The solid red vertical line is the statistics computed from the observed data.

Regression

Robustifying linear regression

- Let us now consider robust regression: the normality assumption of the error terms renders the model quite sensitive to outliers.
- A classic data set in the this literature pertains to hill racing (apparently a somewhat popular sport here in Scotland).
- The data set `hills.txt` contains information on the winning times (in minutes) in 1984 for 35 Scottish hill races, as well as two factors which presumably influence the duration of the race:
  - dist: The distance of the race (in miles).
  - climb: The elevation change (in feet).

Regression

Robustifying linear regression

↪ Fitting a simple linear regression model for time assuming additive linear relationships for dist and climb, we obtain the following residuals

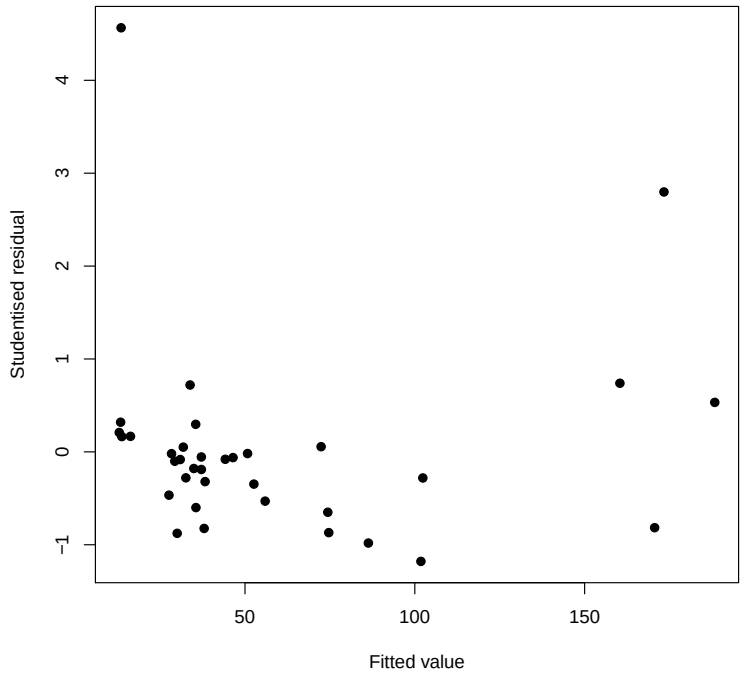

↪ The point at the top left corner, seems suspicious

Regression

Robustifying linear regression

- To reduce the impact of the fit of the model, we can replace the normal distribution with a fat-tailed distribution.
- A natural choice is the *t*-distribution, which can be implemented by simply replacing the normal likelihood with

```
y[i]~dt(mu[i],tau,nu)
```

- Recall that as $\nu \to \infty$, the t-distribution resembles the normal, but for small ν has considerably fatter tails.
- Let us start by trying $\nu = 5$.

Regression

Robustifying linear regression

↪ The bulk of the code is

```
hills=read.table("hills.txt",header=TRUE)
y=hills$time; climb=hills$climb; dist=hills$dist

model_string <- "model{

# Likelihood
for(i in 1:n){
y[i]~dt(mu[i],tau,5)
mu[i]=beta[1]+beta[2]*climb[i]+beta[3]*dist[i]
}

# Prior for beta
for(j in 1:3){
beta[j]~dnorm(mu0,tau0)
}
tau0=1/sigma02

# Prior for the inverse variance
tau~dgamma(a, b)

# Compute the variance
sigma2=1/tau
}"
```

Regression

Robustifying linear regression

↪ Contrasting the results from a normal model versus a t_5 model (posterior mean and standard deviation)

	Climb	Dist
Normal	0.011 (0.0021)	6.204 (0.6150)
t_5	0.0081 (0.0015)	6.565 (0.3089)

↪ There is a modest change in the posterior means (the posterior mean for distance goes up, while the one for climb goes down), and a sizeable drop (roughly 2-fold) in the posterior standard deviation.

Regression

Robustifying linear regression

- ↪ Of course, one may ask, why a t_5 distribution?
- ↪ Since we do not actually know ν, it would be more reasonable to include ν as a parameter in our model; the only condition is that we must place a prior on it.
- ↪ There are several possible options. For instance we could use a similar prior as the one used for the precision.

Regression

Robustifying linear regression

↪ The bulk of the code is now

```
model_string <- "model{
# Likelihood
for(i in 1:n){
y[i]~dt(mu[i],tau,nu)
mu[i]=beta[1]+beta[2]*climb[i]+beta[3]*dist[i]
}

# Prior for beta
for(j in 1:3){
beta[j]~dnorm(mu0,tau0)
}
tau0=1/sigma02

# Prior for the inverse variance
tau~dgamma(a, b)

# Compute the variance
sigma2=1/tau

#Prior for nu
nu~dgamma(c,d)
}"

mu0=0; sigma02=1000; a=0.1; b=0.1; c=0.1; d=0.1
data=list(y=y,climb=climb,dist=dist,n=n,mu0=mu0,sigma02=sigma02,a=a,b=b,c=c,d=d)
```

Regression

Robustifying linear regression

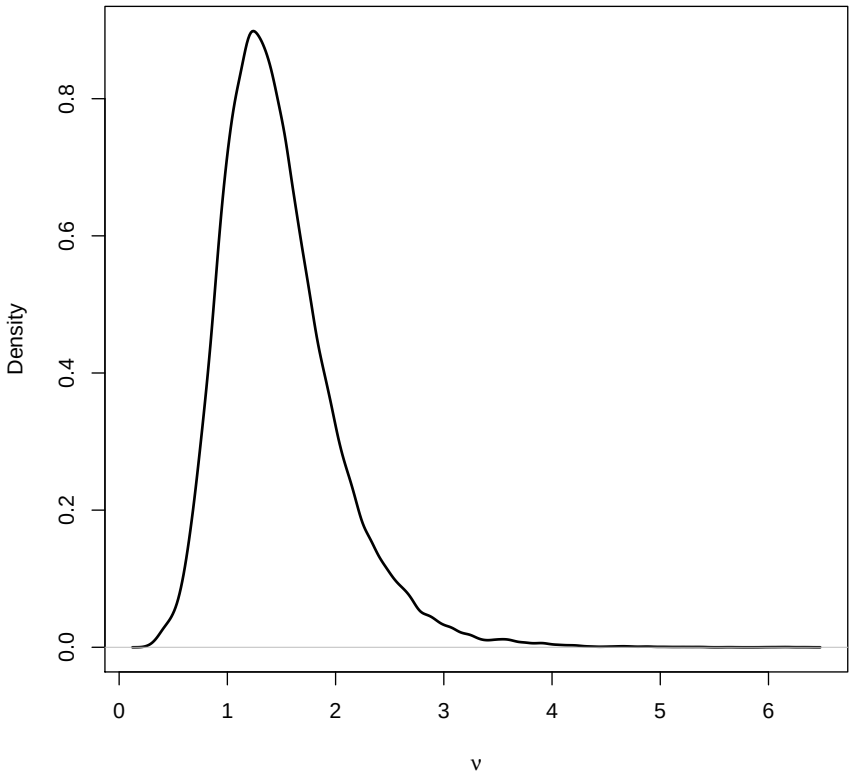

- ↪ The posterior mean (standard deviation) for climb is now 0.0069 (0.0011).
- ↪ The posterior mean (standard deviation) for climb is now 6.543 (0.300).

Regression

Nonlinear regression

↪ Another desirable extension is the ability to fit a nonlinear model.

↪ Carlin and Gelfand (1991) consider data on length (y_i) and age (x_i) measurements for $i = 1, \ldots, n = 27$ dugongs (sea cows) captured off the coast of Queensland.

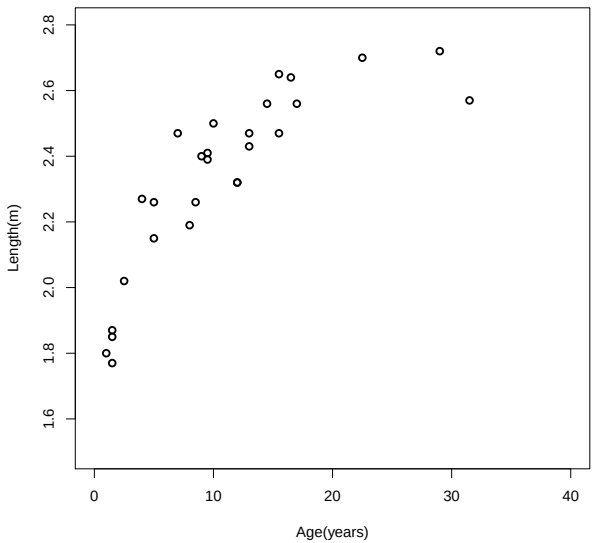

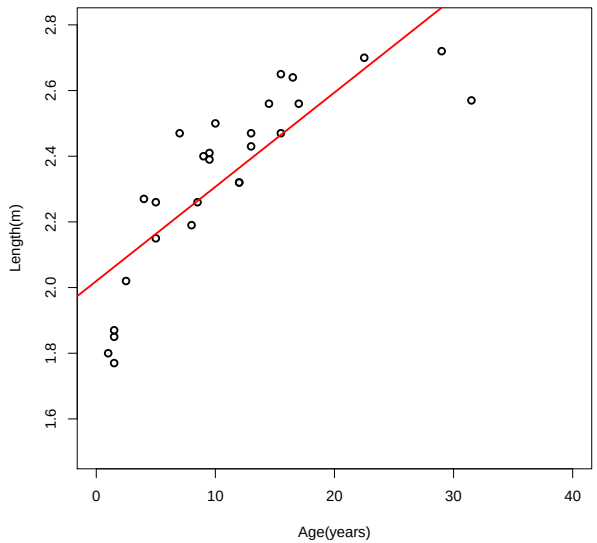

↪ A linear fit to this dataset (shown on the right) is clearly a poor idea.

Regression

Nonlinear regression

- The following nonlinear growth model has been suggested

$$y_i = \alpha - \beta\gamma^{x_i} + \varepsilon_i, \quad i = 1, \ldots, n,$$

where $\alpha > 0$, $\beta > 0$, $0 \leq \gamma \leq 1$ and $\varepsilon_i \overset{\text{iid}}{\sim} \text{N}(0, \sigma^2)$.

- In this model, as explained in Carlin and Louis (2009, p.174), α corresponds to the average length of a fully grown dugong, $(\alpha - \beta)$ is the length of a dugong at birth, and γ determines the growth rate.

- Specifically, lower values of γ produce an initially steep growth curve (rapid progression to adulthood immediately after birth), while higher γ values lead to much more gradual, almost linear growth.

- As stated by these authors, flat priors are suitable for the two "endpoint" parameters α and β, but the harder-to-estimate growth parameter γ benefits from a tighter (albeit uniform) specification.

Regression

Nonlinear regression

↪ The bulk of the code is

```
model_string <- "model{
# Likelihood
for(i in 1:n){
y[i]~dnorm(mu[i],tau)
mu[i]=alpha-beta*pow(gamma,x[i])
}

alpha~dunif(0,10)

beta~dunif(0,10)

gamma~dunif(0.5,1)

# Prior for the inverse variance
tau~dgamma(a, b)

# Compute the variance
sigma2=1/tau
}"

a=0.1; b=0.1
data=list(y=y,x=x,n=n,a=a,b=b)
```

Regression

Nonlinear regression